

worksheet_21

April 19, 2024

1 Worksheet 21

Name: Youxuan Ma

UID: U23330522

1.0.1 Topics

- Logistic Regression
- Gradient Descent

1.1 Logistic Regression

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import sklearn.datasets as datasets
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import PolynomialFeatures

centers = [[0, 0]]
t, _ = datasets.make_blobs(n_samples=750, centers=centers, cluster_std=1,
    random_state=0)

# LINE
def generate_line_data():
    # create some space between the classes
    X = np.array(list(filter(lambda x : x[0] - x[1] < -.5 or x[0] - x[1] > .5,
    t)))
    Y = np.array([1 if x[0] - x[1] >= 0 else 0 for x in X])
    return X, Y

# CIRCLE
def generate_circle_data(t):
    # create some space between the classes
    X = np.array(list(filter(lambda x : (x[0] - centers[0][0])**2 + (x[1] -
    centers[0][1])**2 < 1 or (x[0] - centers[0][0])**2 + (x[1] -
    centers[0][1])**2 > 1.5, t)))
```

```

    Y = np.array([1 if (x[0] - centers[0][0])**2 + (x[1] - centers[0][1])**2 >= 1
↪1 else 0 for x in X])
    return X, Y

# XOR
def generate_xor_data():
    X = np.array([
        [0,0],
        [0,1],
        [1,0],
        [1,1]])
    Y = np.array([x[0]^x[1] for x in X])
    return X, Y

```


a) Using the above code, generate and plot data that is linearly separable.

```

[2]: def plot_data(X, Y):
    plt.scatter(X[:, 0], X[:, 1], c=Y, cmap='bwr', alpha=0.7, edgecolors='k')
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')
    plt.title('Data Scatter Plot')
    plt.show()

X, Y = generate_line_data()
plot_data(X, Y)

```



b) Fit a logistic regression model to the data and print out the coefficients.

```
[3]: model = LogisticRegression().fit(X, Y)
```

```
coefficients = model.coef_  
intercept = model.intercept_  
print("Coefficients:", coefficients)  
print("Intercept:", intercept)
```

```
Coefficients: [[ 4.11128306 -4.10408124]]
```


```
Intercept: [0.06146435]
```

c) Using the coefficients, plot the line through the scatter plot you created in a). (Note: you need to do some math to get the line in the right form)

```
[4]: def plot_decision_boundary(X, Y, model):  
    plt.scatter(X[:, 0], X[:, 1], c=Y, cmap='bwr', alpha=0.7, edgecolors='k')  
    x_values = [np.min(X[:, 0]), np.max(X[:, 0])]  
    y_values = -(model.intercept_ + np.dot(model.coef_[0][0], x_values)) /  
    ↪ model.coef_[0][1]  
    plt.plot(x_values, y_values, label="Decision Boundary", color="black")
```


```
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.title('Logistic Regression Decision Boundary')
plt.legend()
plt.show()

plot_decision_boundary(X, Y, model)
```



d) Using the above code, generate and plot the CIRCLE data.

```
[5]: X_c, Y_c = generate_circle_data(t)
plot_data(X_c, Y_c)
```



e) Notice that the equation of an ellipse is of the form

$$ax^2 + by^2 = c$$

Fit a logistic regression model to an appropriate transformation of X.

```
[6]: X_transformed = np.array([[x[0]**2, x[1]**2] for x in X_c])
model = LogisticRegression().fit(X_transformed, Y_c)

coefficients = model.coef_
intercept = model.intercept_
print("Coefficients:", coefficients)
print("Intercept:", intercept)
```

Coefficients: [[4.91410958 4.97630742]]

Intercept: [-6.45841785]

f) Plot the decision boundary using the code below.

```
[7]: X = X_transformed
Y = Y_c
```

```

# create a mesh to plot in
h = .02 # step size in the mesh
x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
meshData = np.c_[xx.ravel(), yy.ravel()]

fig, ax = plt.subplots()
A = model.predict_proba(meshData)[:, 1].reshape(xx.shape)
Z = model.predict(meshData).reshape(xx.shape)
ax.contourf(xx, yy, A, cmap="RdBu", vmin=0, vmax=1)
ax.axis('off')

# Plot also the training points
ax.scatter(X[:, 0], X[:, 1], c=Y, s=50, alpha=0.9)
plt.show()

```



```

[8]: def plot_decision_boundary(X, Y, model):
    plt.scatter([np.sqrt(x[0]) for x in X], [np.sqrt(x[1]) for x in X], c=Y,
                cmap='viridis', edgecolors='k', alpha=0.7)
    x = np.linspace(0, 1.5, 300)
    y = np.linspace(0, 1.5, 300)

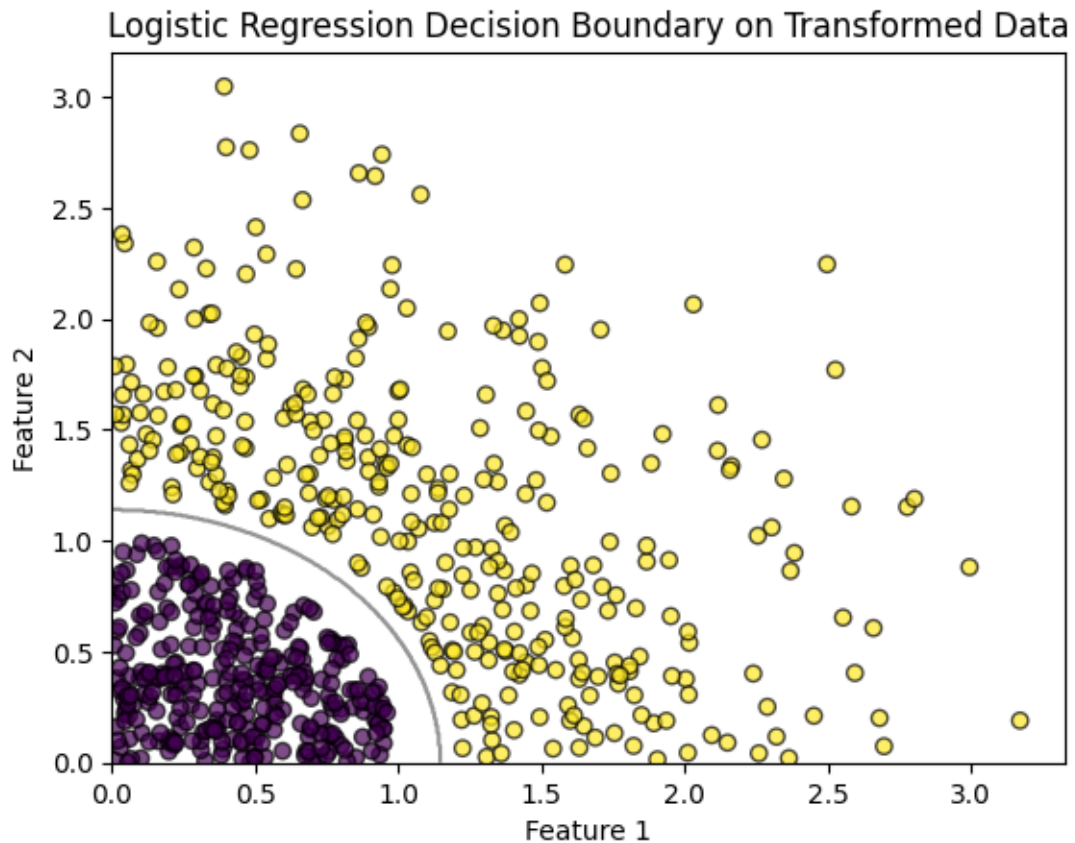
```

```

xx, yy = np.meshgrid(x, y)
zz = model.predict(np.c_[xx.ravel()*2, yy.ravel()*2]).reshape(xx.shape)
plt.contour(xx, yy, zz, levels=[0.5], cmap="Greys", vmin=0, vmax=1)
plt.title('Logistic Regression Decision Boundary on Transformed Data')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()

```

```
plot_decision_boundary(X_transformed, Y_c, model)
```



- g) Plot the XOR data. In this 2D space, the data is not linearly separable, but by introducing a new feature


$$x_3 = x_1 * x_2$$

(called an interaction term) we should be able to find a hyperplane that separates the data in 3D. Plot this new dataset in 3D.

```
[9]: from mpl_toolkits.mplot3d import Axes3D
```

```
X, Y = generate_xor_data()
```

```
ax = plt.axes(projection='3d')
ax.scatter3D(X[:, 0], X[:, 1], X[:, 0]*X[:, 1], c=Y)
plt.show()
```



```
[10]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

# Plot XOR data in 2D
def plot_xor_data(X, Y):
    plt.scatter(X[:, 0], X[:, 1], c=Y, cmap='bwr', edgecolors='k', s=200,
        ↪marker='o')
    plt.xlabel('X1')
    plt.ylabel('X2')
    plt.title('2D Visualization of XOR Data')
    plt.grid(True)
    plt.xticks([0, 1])
    plt.yticks([0, 1])
    plt.show()

# Transform XOR data for 3D visualization
def transform_xor_data(X):
```



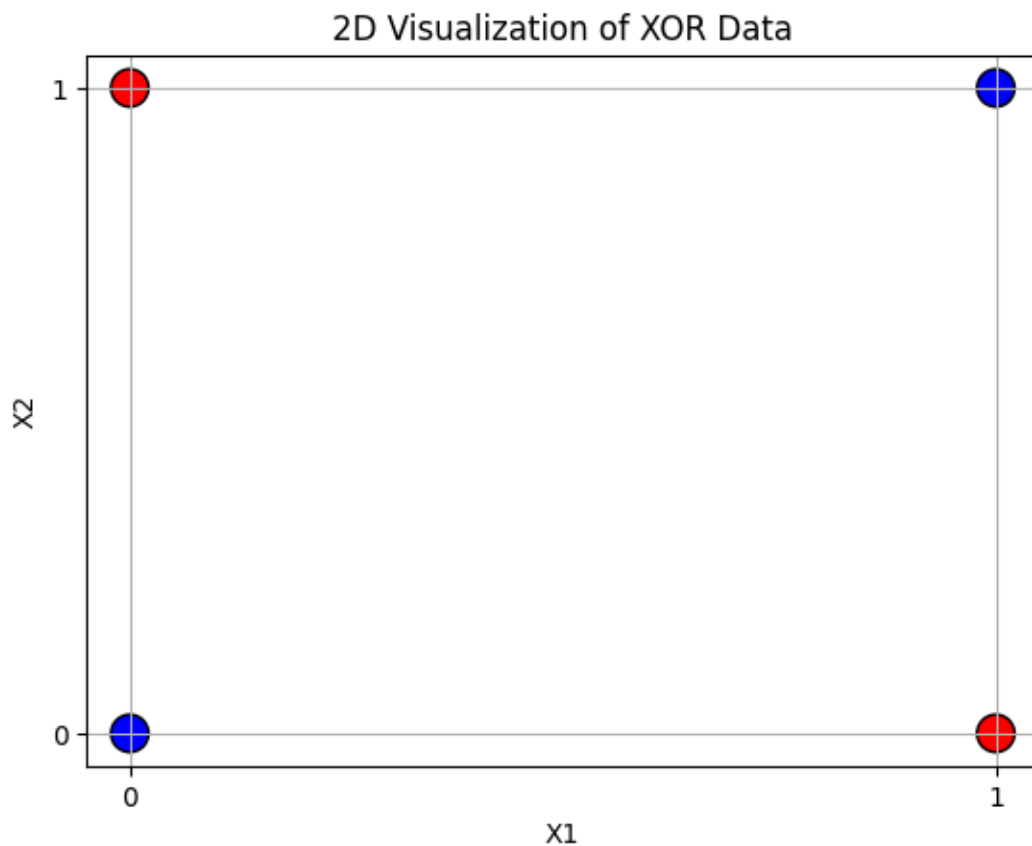
```

X3 = X[:, 0] * X[:, 1]
X_transformed = np.hstack((X, X3[:, np.newaxis]))
return X_transformed


# Plot transformed XOR data in 3D
def plot_3d_xor_data(X, Y):
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')
    ax.scatter(X[:, 0], X[:, 1], X[:, 2], c=Y, cmap='bwr', s=100,
edgecolors='k')
    ax.set_xlabel('X1')
    ax.set_ylabel('X2')
    ax.set_zlabel('X3 = X1*X2')
    ax.set_title('3D Visualization of Transformed XOR Data')
    plt.show()

X, Y = generate_xor_data()
plot_xor_data(X, Y)
X_transformed = transform_xor_data(X)
plot_3d_xor_data(X_transformed, Y)

```



3D Visualization of Transformed XOR Data



h) Apply a logistic regression model using the interaction term. Plot the decision boundary.

```
[11]: X, Y = generate_xor_data()

poly = PolynomialFeatures(interaction_only=True)
lr = LogisticRegression(verbose=0)
model = make_pipeline(poly, lr).fit(X, Y)

# create a mesh to plot in
h = .02 # step size in the mesh
x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
meshData = np.c_[xx.ravel(), yy.ravel()]

fig, ax = plt.subplots()
A = model.predict_proba(meshData)[:, 1].reshape(xx.shape)
Z = model.predict(meshData).reshape(xx.shape)
ax.contourf(xx, yy, A, cmap="RdBu", vmin=0, vmax=1)
```

```
ax.axis('off')

# Plot also the training points
ax.scatter(X[:, 0], X[:, 1], color=Y, s=50, alpha=0.9)
plt.show()
```



```
[12]: X, Y = generate_xor_data()
X_transformed = transform_xor_data(X)
model = LogisticRegression().fit(X_transformed, Y)

coefficients = model.coef_[0]
intercept = model.intercept_[0]
print("Coefficients:", coefficients)
print("Intercept:", intercept)
```

```
Coefficients: [ 0.04285144  0.04285144 -0.43037246]
Intercept: 0.06433821644455287
```

```
[13]: import matplotlib.pyplot as plt
import numpy as np

def plot_2d_decision_boundary(X, Y, model):
    # Set up meshgrid
```

```


x_min, x_max = X[:, 0].min() - 0.1, X[:, 0].max() + 0.1
y_min, y_max = X[:, 1].min() - 0.1, X[:, 1].max() + 0.1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min,
→y_max, 100))

# Calculate z values through the model
Z = model.predict(np.c_[xx.ravel(), yy.ravel(), (xx * yy).ravel()])
Z = Z.reshape(xx.shape)

plt.contourf(xx, yy, Z, alpha=0.4, cmap='bwr')
plt.scatter(X[:, 0], X[:, 1], c=Y, cmap='bwr', edgecolors='k')
plt.title('2D XOR Decision Boundary with Interaction Term')
plt.xlabel('X1')
plt.ylabel('X2')
plt.show()

plot_2d_decision_boundary(X, Y, model)

```



```
[14]: from mpl_toolkits.mplot3d import Axes3D

def plot_3d_decision_boundary(X, Y, model):
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')

    # Set up meshgrid
    x_min, x_max = X[:, 0].min() - 0.1, X[:, 0].max() + 0.1
    y_min, y_max = X[:, 1].min() - 0.1, X[:, 1].max() + 0.1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 50), np.linspace(y_min,
↪y_max, 50))
    zz = model.predict(np.c_[xx.ravel(), yy.ravel(), (xx * yy).ravel()])


    # Plot the surface
    zz = zz.reshape(xx.shape)
    ax.plot_surface(xx, yy, zz, color='b', alpha=0.3)

    # Scatter plot of original data points
    ax.scatter(X[:, 0], X[:, 1], X[:, 2], c=Y, cmap='bwr', edgecolors='k',
↪depthshade=False)
    ax.set_xlabel('X1')
    ax.set_ylabel('X2')
    ax.set_zlabel('X3 = X1*X2')
    ax.set_title('3D XOR Decision Boundary with Interaction Term')

    plt.show()

plot_3d_decision_boundary(X_transformed, Y, model)
```

3D XOR Decision Boundary with Interaction Term



```
[15]: from mpl_toolkits.mplot3d import Axes3D
X, Y = generate_xor_data()

for i in range(2000):
    for solver in ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag',
                  ↪ 'saga']:
        X_transform = PolynomialFeatures(interaction_only=True,
        ↪ include_bias=False).fit_transform(X)
        model = LogisticRegression(verbose=0, solver=solver, random_state=i,
        ↪ max_iter=10000)
        model.fit(X_transform, Y)

        if model.score(X_transform, Y) > .75:
            print("random state = ", i)
            print("solver = ", solver)
            print(model.score(X_transform, Y))
            break

print(model.coef_)
print(model.intercept_)
```

```

xx, yy = np.meshgrid([x / 10 for x in range(-1, 11)], [x / 10 for x in
    ↪range(-1, 11)])
z = - model.intercept_ / model.coef_[0][2] - model.coef_[0][0] * xx / model.
    ↪coef_[0][2] - model.coef_[0][1] * yy / model.coef_[0][2]


ax = plt.axes(projection='3d')
ax.scatter3D(X[:, 0], X[:, 1], X[:, 0] * X[:, 1], c=Y)
ax.plot_surface(xx, yy, z, alpha=0.5)
plt.show()

```

```

[[ 0.04261388  0.04263475 -0.43040403]]
[0.06489582]

```



- i) Using the code below that generates 3 concentric circles, fit a logistic regression model to it and plot the decision boundary.

```

[16]: t, _ = datasets.make_blobs(n_samples=1500, centers=centers, cluster_std=2,
    random_state=0)

# CIRCLES
def generate_circles_data(t):
    def label(x):
        if x[0]**2 + x[1]**2 >= 2 and x[0]**2 + x[1]**2 < 8:

```

```

        return 1
    if x[0]**2 + x[1]**2 >= 8:
        return 2
    return 0
# create some space between the classes
X = np.array(list(filter(lambda x : (x[0]**2 + x[1]**2 < 1.8 or x[0]**2 +
x[1]**2 > 2.2) and (x[0]**2 + x[1]**2 < 7.8 or x[0]**2 + x[1]**2 > 8.2), t)))
Y = np.array([label(x) for x in X])
return X, Y

X, Y = generate_circles_data(t)

poly = PolynomialFeatures(2)
lr = LogisticRegression(verbose=2)
model = make_pipeline(poly, lr).fit(X, Y)

def plot_decision_boundary(X, Y, model):
    # Set up the mesh grid
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 500),
                          np.linspace(y_min, y_max, 500))

    # Predict over the grid
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.figure(figsize=(10, 8))
    plt.contourf(xx, yy, Z, alpha=0.5, cmap='viridis')
    plt.scatter(X[:, 0], X[:, 1], c=Y, s=20, edgecolor='k', cmap='viridis')
    plt.title('Decision Boundary for Concentric Circles')
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')
    plt.show()

plot_decision_boundary(X, Y, model)

```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 21 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.09861D+00 |proj g|= 1.56782D+00

At iterate	1	f=	8.95383D-01	proj g =	5.58383D-01
At iterate	2	f=	8.27195D-01	proj g =	3.32189D-01
At iterate	3	f=	7.96120D-01	proj g =	6.06694D-01
At iterate	4	f=	7.66577D-01	proj g =	1.62504D-01
At iterate	5	f=	7.38227D-01	proj g =	1.74906D-01
At iterate	6	f=	6.22098D-01	proj g =	1.59122D-01
At iterate	7	f=	4.78311D-01	proj g =	1.20393D-01
At iterate	8	f=	2.88632D-01	proj g =	2.29261D-01
At iterate	9	f=	2.73371D-01	proj g =	3.32278D-01
At iterate	10	f=	1.61027D-01	proj g =	1.86164D-01
At iterate	11	f=	1.31057D-01	proj g =	8.20086D-02
At iterate	12	f=	1.09857D-01	proj g =	4.31233D-02
At iterate	13	f=	9.54732D-02	proj g =	3.14860D-02
At iterate	14	f=	8.82113D-02	proj g =	3.19970D-02
At iterate	15	f=	8.65265D-02	proj g =	2.58420D-02
At iterate	16	f=	8.51986D-02	proj g =	7.71859D-03
At iterate	17	f=	8.49323D-02	proj g =	4.32702D-03
At iterate	18	f=	8.47892D-02	proj g =	2.54731D-03
At iterate	19	f=	8.45880D-02	proj g =	4.72692D-03
At iterate	20	f=	8.43495D-02	proj g =	9.76915D-03
At iterate	21	f=	8.39228D-02	proj g =	1.41123D-02
At iterate	22	f=	8.30751D-02	proj g =	1.33183D-02
At iterate	23	f=	8.27346D-02	proj g =	1.95171D-02
At iterate	24	f=	8.16237D-02	proj g =	1.19741D-02

At iterate	25	f=	7.96364D-02	proj g =	7.19066D-03
At iterate	26	f=	7.76664D-02	proj g =	1.85242D-02
At iterate	27	f=	7.46684D-02	proj g =	3.03654D-02
At iterate	28	f=	7.05634D-02	proj g =	5.30341D-02
At iterate	29	f=	6.54279D-02	proj g =	3.59637D-02
At iterate	30	f=	6.18421D-02	proj g =	2.28137D-02
At iterate	31	f=	5.85717D-02	proj g =	6.21250D-03
At iterate	32	f=	5.74600D-02	proj g =	5.64093D-03
At iterate	33	f=	5.58955D-02	proj g =	1.31517D-02
At iterate	34	f=	5.38386D-02	proj g =	1.53870D-02
At iterate	35	f=	5.32192D-02	proj g =	1.92596D-02
At iterate	36	f=	5.19024D-02	proj g =	3.39938D-03
At iterate	37	f=	5.17519D-02	proj g =	1.98737D-03
At iterate	38	f=	5.16150D-02	proj g =	1.95118D-03
At iterate	39	f=	5.15147D-02	proj g =	2.96761D-03
At iterate	40	f=	5.14626D-02	proj g =	2.59971D-03
At iterate	41	f=	5.13984D-02	proj g =	1.62724D-03
At iterate	42	f=	5.13370D-02	proj g =	9.99570D-04
At iterate	43	f=	5.12664D-02	proj g =	1.37624D-03
At iterate	44	f=	5.11763D-02	proj g =	1.60635D-03
At iterate	45	f=	5.10465D-02	proj g =	1.34433D-03
At iterate	46	f=	5.10047D-02	proj g =	4.95510D-03
At iterate	47	f=	5.08214D-02	proj g =	4.46364D-03
At iterate	48	f=	5.07167D-02	proj g =	1.08309D-03

At iterate	49	f=	5.06903D-02	proj g =	8.37843D-04
At iterate	50	f=	5.06685D-02	proj g =	4.53040D-04
At iterate	51	f=	5.06555D-02	proj g =	1.17846D-03
At iterate	52	f=	5.06408D-02	proj g =	7.58418D-04
At iterate	53	f=	5.06250D-02	proj g =	9.22698D-04
At iterate	54	f=	5.05228D-02	proj g =	1.71929D-03
At iterate	55	f=	5.03782D-02	proj g =	1.97601D-03
At iterate	56	f=	5.02844D-02	proj g =	2.05844D-03
At iterate	57	f=	5.01694D-02	proj g =	9.05956D-04
At iterate	58	f=	5.01232D-02	proj g =	7.69174D-04
At iterate	59	f=	5.01149D-02	proj g =	8.32515D-04
At iterate	60	f=	5.01087D-02	proj g =	1.84661D-04
At iterate	61	f=	5.01070D-02	proj g =	1.57103D-04
At iterate	62	f=	5.01059D-02	proj g =	4.65643D-04
At iterate	63	f=	5.01039D-02	proj g =	2.51792D-04
At iterate	64	f=	5.01027D-02	proj g =	1.24677D-04
At iterate	65	f=	5.01016D-02	proj g =	1.21590D-04
At iterate	66	f=	5.01008D-02	proj g =	9.65235D-05

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value


* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
21	66	75	1	0	0	9.652D-05	5.010D-02

F = 5.0100827208558775E-002

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

This problem is unconstrained.



1.2 Gradient Descent

Recall in Linear Regression we are trying to find the line

$$y = X\beta$$

that minimizes the sum of square distances between the predicted y and the y we observed in our dataset:

$$\mathcal{L}(\beta) = \|\mathbf{y} - X\beta\|^2$$


We were able to find a global minimum to this loss function but we will try to apply gradient descent to find that same solution.

a) Implement the loss function to complete the code and plot the loss as a function of beta.

```
[17]: from mpl_toolkits import mplot3d
import numpy as np
import matplotlib.pyplot as plt

beta = np.array([ 1 , .5 ])
xlin = -10.0 + 20.0 * np.random.random(100)
X = np.column_stack([np.ones((len(xlin), 1)), xlin])
y = beta[0]+(beta[1]*xlin)+np.random.randn(100)

fig, ax = plt.subplots()
ax.plot(xlin, y, 'ro', markersize=4)
ax.set_xlim(-10, 10)
ax.set_ylim(min(y), max(y))
plt.show()
```



```

[18]: b0 = np.arange(-5, 4, 0.1)
      b1 = np.arange(-5, 4, 0.1)
      b0, b1 = np.meshgrid(b0, b1)

      def loss(X, y, beta):
          residuals = y - X @ beta
          return np.dot(residuals, residuals)

      def get_cost(B0, B1):
          res = []
          for b0, b1 in zip(B0, B1):
              line = []
              for i in range(len(b0)):
                  beta = np.array([b0[i], b1[i]])
                  line.append(loss(X, y, beta))
              res.append(line)
          return np.array(res)


      cost = get_cost(b0, b1)

      # Creating figure
      fig = plt.figure(figsize=(14, 9))
      ax = plt.axes(projection='3d')
      ax.set_xlim(-6, 4)
      ax.set_xlabel(r'$\beta_0$')
      ax.set_ylabel(r'$\beta_1$')
      ax.set_ylim(-5, 4)
      ax.set_zlim(0, 30000)

      # Creating plot
      ax.plot_surface(b0, b1, cost, alpha=.7)

      # show plot
      plt.show()

```



Since the loss is

$$\mathcal{L}(\beta) = \|\mathbf{y} - X\beta\|^2 = \beta^T X^T X \beta - 2^T X^T \mathbf{y} + \mathbf{y}^T \mathbf{y}$$

The gradient is

$$\nabla_{\beta} \mathcal{L}(\beta) = 2X^T X \beta - 2X^T \mathbf{y}$$

b) Implement the gradient function below and complete the gradient descent algorithm

```
[19]: import numpy as np
      from PIL import Image as im
      import matplotlib.pyplot as plt
```

```

TEMPFILE = "temp.png"

def snap(betas, losses):
    fig = plt.figure(figsize=(14, 9))
    ax = plt.axes(projection='3d')
    ax.view_init(20, -20)
    ax.set_xlim(-5, 4)
    ax.set_xlabel(r'$\beta_0$')
    ax.set_ylabel(r'$\beta_1$')
    ax.set_ylim(-5, 4)
    ax.set_zlim(0, 30000)

    ax.plot_surface(b0, b1, cost, color='b', alpha=.7)
    ax.plot(np.array(betas)[: ,0], np.array(betas)[: ,1], losses, 'o-', c='r',
    ↪markersize=10, zorder=10)
    fig.canvas.draw()

    img = np.frombuffer(fig.canvas.buffer_rgba(), dtype=np.uint8)
    img = img.reshape(fig.canvas.get_width_height()[::-1] + (4,))
    plt.close(fig)

    img = img[... , :3]
    return im.fromarray(img)

def gradient(X, y, beta):
    return 2 * X.T @ X @ beta - 2 * X.T @ y

def gradient_descent(X, y, beta_hat, learning_rate, epochs, images):
    losses = [loss(X, y, beta_hat)]
    betas = [beta_hat]

    for _ in range(epochs):
        images.append(snap(betas, losses))
        grad = gradient(X, y, beta_hat)
        beta_hat = beta_hat - learning_rate * grad

        losses.append(loss(X, y, beta_hat))
        betas.append(beta_hat)

    return np.array(betas), np.array(losses)

beta_start = np.array([-5, -2])
learning_rate = 0.0002 # try .0005

```



```

epochs = 50
images = []
betas, losses = gradient_descent(X, y, beta_start, learning_rate, epochs,
    ↪images)

images[0].save(
    'gd.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=500
)

```

c) Use the code above to create an animation of the linear model learned at every epoch.

```

[20]: def snap_model(beta):
    xplot = np.linspace(-10,10,50)
    yestplot = beta[0] + beta[1] * xplot
    fig, ax = plt.subplots()
    ax.plot(xplot, yestplot, 'b-', lw=2)
    ax.plot(xlin, y, 'ro', markersize=4)
    ax.set_xlim(-10, 10)
    ax.set_ylim(min(y), max(y))
    fig.savefig(TEMPFILE)
    plt.close()
    return im.fromarray(np.asarray(im.open(TEMPFILE)))

def gradient_descent(X, y, beta_hat, learning_rate, epochs, images):
    losses = [loss(X, y, beta_hat)]
    betas = [beta_hat]

    for _ in range(epochs):
        images.append(snap_model(beta_hat))
        beta_hat = beta_hat - learning_rate * gradient(X, y, beta_hat)

        losses.append(loss(X, y, beta_hat))
        betas.append(beta_hat)

    return np.array(betas), np.array(losses)

images = []
betas, losses = gradient_descent(X, y, beta_start, learning_rate, 100, images)

images[0].save(

```

```

    'model.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=200
)

```

In logistic regression, the `loss` is the negative log-likelihood

$$l(\beta) = -\frac{1}{N} \sum_{i=1}^N y_i \log(\sigma(x_i \beta)) + (1 - y_i) \log(1 - \sigma(x_i \beta))$$

the gradient of which is:

$$\nabla_{\beta} l(\beta) = \frac{1}{N} \sum_{i=1}^N x_i (y_i - \sigma(x_i \beta))$$

d) Plot the loss as a function of β .

```

[21]: from mpl_toolkits import mplot3d
import numpy as np
import matplotlib.pyplot as plt
import sklearn.datasets as datasets


centers = [[0, 0]]
t, _ = datasets.make_blobs(n_samples=100, centers=centers, cluster_std=2,
    random_state=0)

# LINE
def generate_line_data():
    # create some space between the classes
    X = t
    Y = np.array([1 if x[0] - x[1] >= 0 else 0 for x in X])
    return X, Y

X, y = generate_line_data()

cs = np.array([x for x in 'gb'])
fig, ax = plt.subplots()
ax.scatter(X[:, 0], X[:, 1], color=cs[y].tolist(), s=50, alpha=0.9)
plt.show()

```



```
[22]: b0 = np.arange(-20, 20, 0.1)
      b1 = np.arange(-20, 20, 0.1)
      b0, b1 = np.meshgrid(b0, b1)

      def sigmoid(x):
          e = np.exp(x)
          return e / (1 + e)

      def loss(X, y, beta):
          predictions = sigmoid(X @ beta)
          epsilon = 1e-9 # Adding a small constant to prevent log(0)
          loss = -np.mean(y * np.log(predictions + epsilon) + (1 - y) * np.log(1 -
          ↪ predictions + epsilon))
          return loss

      def get_cost(B0, B1):
          res = []
          for b0, b1 in zip(B0, B1):
              line = []
              for i in range(len(b0)):
```

```

        beta = np.array([b0[i], b1[i]])
        line.append(loss(X, y, beta))
    res.append(line)
    return np.array(res)


cost = get_cost(b0, b1)

# Creating figure
fig = plt.figure(figsize=(14, 9))
ax = plt.axes(projection='3d')
ax.set_xlim(-20, 20)
ax.set_xlabel(r'$\beta_0$')
ax.set_ylabel(r'$\beta_1$')
ax.set_ylim(-20, 20)
ax.set_zlim(0, 10)

# Creating plot
ax.plot_surface(b0, b1, cost, alpha=.7)

# show plot
plt.show()

```



e) Plot the loss at each iteration of the gradient descent algorithm.

```
[23]: import numpy as np
from PIL import Image as im
import matplotlib.pyplot as plt

TEMPFILE = "temp.png"

def snap(betas, losses):
    fig = plt.figure(figsize=(14, 9))
    ax = plt.axes(projection='3d')
    ax.view_init(10, 10)
    ax.set_xlim(-20, 20)
```

```

ax.set_xlabel(r'$\beta_0$')
ax.set_ylabel(r'$\beta_1$')
ax.set_ylim(-20, 20)
ax.set_zlim(0, 10)

ax.plot_surface(b0, b1, cost, color='b', alpha=.7)
ax.plot(np.array(betas)[: , 0], np.array(betas)[: , 1], losses, 'o-', c='r',
↪markersize=10, zorder=10)
plt.close()

fig.canvas.draw()
img = np.frombuffer(fig.canvas.buffer_rgba(), dtype=np.uint8)
img = img.reshape(fig.canvas.get_width_height()[::-1] + (4,))
img = img[..., :3]

return im.fromarray(img)

def gradient(X, y, beta):
    predictions = sigmoid(X @ beta)
    errors = y - predictions
    grad = -X.T @ errors / len(y)
    return grad

def gradient_descent(X, y, beta_hat, learning_rate, epochs, images):
    losses = [loss(X, y, beta_hat)]
    betas = [beta_hat]

    for _ in range(epochs):
        images.append(snap(betas, losses))
        beta_hat = beta_hat - learning_rate * gradient(X, y, beta_hat)

        losses.append(loss(X, y, beta_hat))
        betas.append(beta_hat)

    return np.array(betas), np.array(losses)

beta_start = np.array([-5, -2])
learning_rate = 0.1
images = []
betas, losses = gradient_descent(X, y, beta_start, learning_rate, 50, images)

images[0].save(
    'gd_logit.gif',

```

```

optimize=False,
save_all=True,
append_images=images[1:],
loop=0,
duration=500
)

```

f) Create an animation of the logistic regression fit at every epoch.

```

[24]: def snap_fit(beta, X, y, epoch, ax=None):
    if ax is None:
        fig, ax = plt.subplots()
        x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
        y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
        xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min,
        ↪ y_max, 100))

        Z = sigmoid(np.c_[xx.ravel(), yy.ravel()] @ beta)
        Z = Z.reshape(xx.shape)

        cs = np.array([x for x in 'gb'])
        ax.contourf(xx, yy, Z, levels=[0, 0.5, 1], cmap='RdBu', alpha=0.5)
        ax.scatter(X[:, 0], X[:, 1], color=cs[y].tolist(), s=50, alpha=0.9,
        ↪ edgecolors='k')

        ax.set_xlim(xx.min(), xx.max())
        ax.set_ylim(yy.min(), yy.max())
        ax.set_title(f'Epoch {epoch}')
        plt.close()

        fig.canvas.draw()
        img = np.frombuffer(fig.canvas.buffer_rgba(), dtype=np.uint8)
        img = img.reshape(fig.canvas.get_width_height()[::-1] + (4,))
        img = img[..., :3]

        return im.fromarray(img)

def gradient_descent(X, y, beta_hat, learning_rate, epochs, images):
    losses = [loss(X, y, beta_hat)]
    betas = [beta_hat]
    images.append(snap_fit(beta_hat, X, y, 0))

    for epoch in range(1, epochs + 1):
        grad = gradient(X, y, beta_hat)
        beta_hat = beta_hat - learning_rate * grad
        current_loss = loss(X, y, beta_hat)

```

```

        losses.append(current_loss)
        betas.append(beta_hat)
        images.append(snap_fit(beta_hat, X, y, epoch))

    return np.array(betas), np.array(losses)

beta_start = np.array([-3, -2])
learning_rate = 0.1
epochs = 100
images = []

betas, losses = gradient_descent(X, y, beta_start, learning_rate, epochs,
    ↪ images)

images[0].save(
    'logit_model_fit.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=200
)

```

g) Modify the above code to evaluate the gradient on a random batch of the data. Overlay the true loss curve and the approximation of the loss in your animation.

```

[25]: def get_batch(X, y, batch_size):
        idx = np.random.randint(0, len(y), batch_size)
        return X[idx], y[idx]

    def gradient(X, y, beta):
        predictions = sigmoid(X @ beta)
        errors = y - predictions
        grad = -X.T @ errors / len(y)
        return grad

    def gradient_descent(X, y, beta_hat, learning_rate, epochs, batch_size, images,
    ↪ loss_data):
        for epoch in range(epochs):
            X_batch, y_batch = get_batch(X, y, batch_size)
            grad = gradient(X_batch, y_batch, beta_hat)
            beta_hat = beta_hat - learning_rate * grad

            approx_loss = loss(X_batch, y_batch, beta_hat)
            true_loss = loss(X, y, beta_hat)

            loss_data.append((true_loss, approx_loss))

```



```

        images.append(snap_fit(beta_hat, X, y, epoch, loss_data))

    return beta_hat

def snap_fit(beta, X, y, epoch, loss_data, ax=None):
    if ax is None:
        fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 7))
    else:
        ax1, ax2 = ax

    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min,
↪y_max, 100))

    Z = sigmoid(np.c_[xx.ravel(), yy.ravel()] @ beta)
    Z = Z.reshape(xx.shape)
    cs = np.array([x for x in 'gb'])
    ax1.contourf(xx, yy, Z, levels=[0, 0.5, 1], cmap='RdBu', alpha=0.5)
    ax1.scatter(X[:, 0], X[:, 1], color=cs[y].tolist(), s=50, alpha=0.9,
↪edgecolors='k')

    ax1.set_xlim(xx.min(), xx.max())
    ax1.set_ylim(yy.min(), yy.max())
    ax1.set_title(f'Epoch {epoch}')

    true_losses, approx_losses = zip(*loss_data)
    ax2.plot(true_losses, label='True Loss')
    ax2.plot(approx_losses, label='Approximated Loss')
    ax2.legend()
    ax2.set_title('Loss over epochs')
    ax2.set_xlabel('Epoch')
    ax2.set_ylabel('Loss')

    plt.close(fig)
    fig.canvas.draw()
    img = np.frombuffer(fig.canvas.buffer_rgba(), dtype=np.uint8)
    img = img.reshape(fig.canvas.get_width_height()[::-1] + (4,))
    img = img[..., :3]

    return im.fromarray(img)

beta_start = np.array([-3, -2])
learning_rate = 0.1
epochs = 100
batch_size = 20
images = []

```

```

loss_data = []

final_beta = gradient_descent(X, y, beta_start, learning_rate, epochs,
    ↪batch_size, images, loss_data)

images[0].save(
    'batch_logit_fit.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=200
)

```

h) Below is a sandox where you can get intuition about how to tune gradient descent parameters:

```

[26]: import numpy as np
from PIL import Image as im
import matplotlib.pyplot as plt

TEMPFILE = "temp.png"

def snap(x, y, pts, losses, grad):
    fig = plt.figure(figsize=(14, 9))
    ax = plt.axes(projection='3d')
    ax.view_init(20, -20)
    ax.plot_surface(x, y, loss(np.array([x, y])), color='r', alpha=.4)
    ax.plot(np.array(pts)[: ,0], np.array(pts)[: ,1], losses, 'o-', c='b',
    ↪markersize=10, zorder=10)
    ax.plot(np.array(pts)[-1,0], np.array(pts)[-1,1], -1, 'o-', c='b', alpha=.
    ↪5, markersize=7, zorder=10)

    # Plot Gradient Vector
    X, Y, Z = [pts[-1][0]], [pts[-1][1]], [-1]
    U, V, W = [-grad[0]], [-grad[1]], [0]
    ax.quiver(X, Y, Z, U, V, W, color='g')
    fig.savefig(TEMPFILE)
    plt.close()
    return im.fromarray(np.asarray(im.open(TEMPFILE)))

def loss(x):
    return np.sin(sum(x**2)) # change this

def gradient(x):
    return 2 * x * np.cos(sum(x**2)) # change this

def gradient_descent(x, y, init, learning_rate, epochs):

```

```

images, losses, pts = [], [loss(init)], [init]
for _ in range(epochs):
    grad = gradient(init)
    images.append(snap(x, y, pts, losses, grad))
    init = init - learning_rate * grad
    losses.append(loss(init))
    pts.append(init)
return images

init = np.array([-0.5, -0.5]) # change this
learning_rate = 1.394 # change this
x, y = np.meshgrid(np.arange(-2, 2, 0.1), np.arange(-2, 2, 0.1)) # change this
images = gradient_descent(x, y, init, learning_rate, 12)

images[0].save(
    'gradient_descent.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=500
)

```